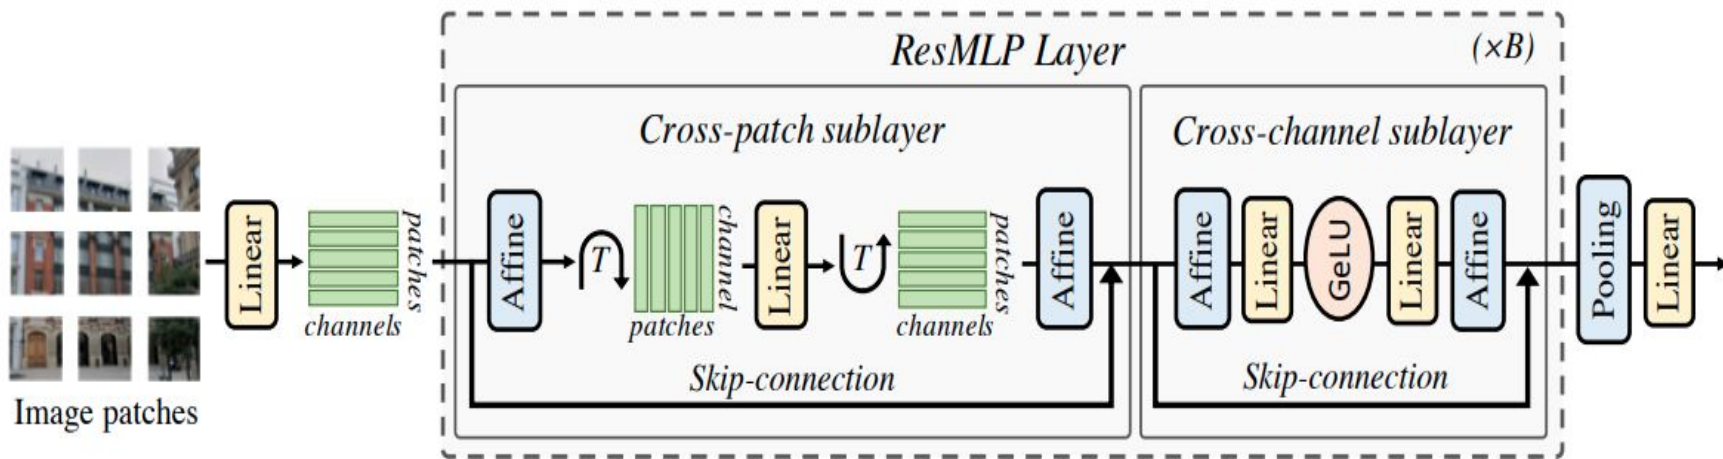




**RESMLP**

**Implementation**

# Architecture



# Hyperparameters

Dataset = Cifar10

Image size = 32

patch\_size = 16

num\_layers = 12

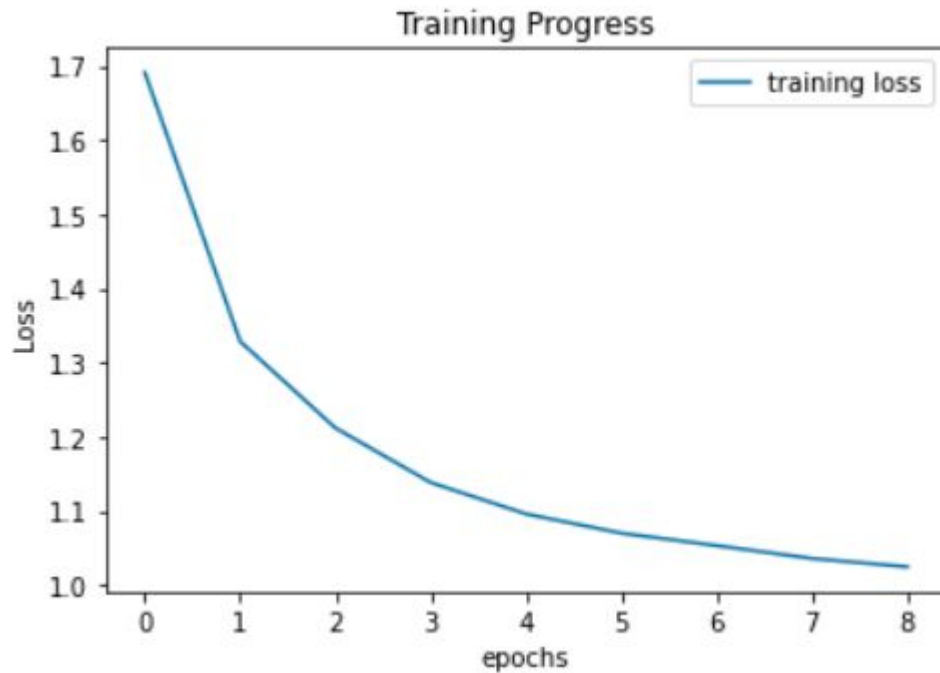
batch\_size = 32

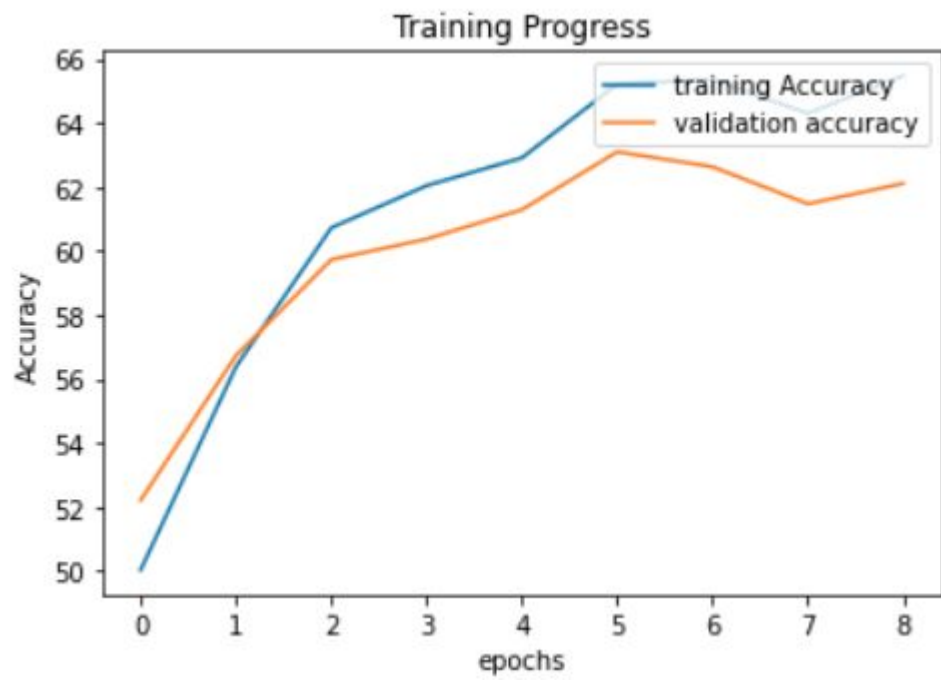
lr =  $1e-2$

epochs = 10

Number of parameters = 14,978,170

# Progress so far





# Training remarks

- The architecture is straightforward
- Training is fairly stable with smooth curves

**Self supervised with Dino**

# Demo



